

# Classification - Assignment 7

## Data and Package Import

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import pylab as plt
```

```

In [2]: from sklearn.datasets import make_blobs, make_moons, make_circles
np.random.seed(4)

noisiness = 1

X_blob, y_blob = make_blobs(n_samples = 200, centers = 2, cluster_std = 2 * noisiness, n_feat
X_mc, y_mc = make_blobs(n_samples = 200, centers = 3, cluster_std = 0.5 * noisiness, n_feat
X_circles, y_circles = make_circles(n_samples = 200, factor = 0.3, noise = 0.1 * noisiness)
X_moons, y_moons = make_moons(n_samples = 200, noise = 0.25 * noisiness)

N_include = 30
idxs = []
Ni = 0
for i, yi in enumerate(y_moons):
    if yi == 1 and Ni < N_include:
        idxs.append(i)
        Ni += 1
    elif yi == 0:
        idxs.append(i)

y_moons = y_moons[idxs]
X_moons = X_moons[idxs]


fig, axes = plt.subplots(1, 4, figsize = (15, 3), dpi = 200)

all_datasets = [[X_blob, y_blob], [X_mc, y_mc], [X_circles, y_circles], [X_moons, y_moons]]

labels = ['Dataset 1', 'Dataset 2', 'Dataset 3', 'Dataset 4']
for i, Xy_i in enumerate(all_datasets):
    Xi, yi = Xy_i
    axes[i].scatter(Xi[:, 0], Xi[:, 1], c = yi)
    axes[i].set_title(labels[i])
    axes[i].set_xlabel('$x_0$')
    axes[i].set_ylabel('$x_1$')

fig.subplots_adjust(wspace = 0.4);

```



```

In [3]: df = pd.read_csv('data/perovskite_data.csv')
X_perov = df[['nA', 'nB', 'nX', 'rA (Ang)', 'rB (Ang)', 'rX (Ang)', 't', 'tau']].values
y_perov = df['exp_label'].values

```

## 1. k-nearest Neighbors Model

## 1-NN

**Calculate the accuracy of a 1-nearest Neighbors model for the training data.**

A 1-nearest Neighbors model considers a point as its own nearest neighbors.

Accuracy will be just 1. 1-NN simply memorizes the label of each data point.

**Will this be a reliable indicator of its accuracy for testing data?**

Briefly explain your answer.

No. Even though the accuracy is 1, this cannot be a reliable indicator, since we know that this score is due to a simple memorization.

## Weighted Neighbors Classification

Instead of selecting the k-nearest neighbors to vote, we could design an algorithm where all neighbors get to vote, but their vote is weighted inversely to their distance from the point of interest:

$$y_i = \sum_j y_j / (||x_i - x_j||)$$

where  $j$  is an index over all training points.

The class will be assigned as follows:

- class 1 if  $y_i \geq 0$
- class -1 if  $y_i < 0$

```
In [4]: def distance(x1, x2):  
        return np.linalg.norm(x1 - x2, 2)
```

```
In [5]: def get_neighbor(x, x_list):  
        dist_pairs = []  
        for i, xi in enumerate(x_list):  
            dist = distance(x, xi)  
            dist_pairs.append([dist, i])  
        return dist_pairs
```

**Write a function that assigns a class to a point.**

The function should take the followings as arguments:

- a single point `x`
- a list of training points `x_list`
- a list of training labels `y_list`

You may want to use functions above. You will also need to add a statement to avoid dividing by zero if the point is in the training set. If the distance between 2 points is zero, then the label from the same point in the training set should be used (e.g. if

```
In [6]: def assign_class(x, x_list, y_list):
        neighbors = get_neighbor(x, x_list)
        vote = 0

        for i, xi in enumerate(neighbors):
            if neighbors[i][0] == 0:
                vote = y_list[i]
                break
            vote += y_list[i] / neighbors[i][0]

        if vote >= 0:
            assignment = 1
        else:
            assignment = -1

        return assignment
```

**Write a function that returns the prediction for a given list of testing points.**

The function should take the followings as arguments:

- a list of testing points `X`
- a list of training points `X_train`
- a list of training labels `y_train`

```
In [7]: def weighted_neighbors(X, X_train, y_train):
        y_hat = []
        for x in X:
            assign = assign_class(x, X_train, y_train)
            y_hat.append(assign)
        return y_hat
```

**Train the model for the perovskite dataset using a random selection of 75% of the data as training data.**

```
In [8]: from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(X_perov, y_perov, test_size = 0.25)
        y_predict = weighted_neighbors(X_perov, X_train, y_train)
```

**Compute the accuracy and precision of the prediction.**

```
In [9]: from sklearn.metrics import accuracy_score, precision_score
```

```
acc = accuracy_score(y_perov, y_predict)
prec = precision_score(y_perov, y_predict)
```

```
print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(prec))
```

```
Accuracy: 0.9635416666666666
Precision: 0.9506172839506173
```

**Train a 5-NN model using the same training data.**

```
In [10]: from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors = 5)
```

```
knn.fit(X_train, y_train)
y_predict = knn.predict(X_perov)
```

**Compute the accuracy and precision.**

```
In [11]: acc = accuracy_score(y_perov, y_predict)
prec = precision_score(y_perov, y_predict)
```

```
print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(prec))
```

```
Accuracy: 0.9322916666666666
Precision: 0.9391025641025641
```

## 2. Multi-dimensional Classification

### Simple logistic regression

**Train a logistic regression model using all columns except the `tau` column of the perovskite dataset.**

You may use some functions that have been already built in the previous assignments.

```
In [12]: from sklearn.linear_model import LogisticRegression
```

```
X = X_perov[:, :-1]
```


```
logreg = LogisticRegression(penalty = 'none') # no regularization
logreg.fit(X, y_perov);
```

**Plot the confusion matrix.**

```
In [13]: from sklearn.metrics import confusion_matrix
import seaborn as sns

y_predict = logreg.predict(X)
cm = confusion_matrix(y_perov, y_predict)

fig, ax = plt.subplots(figsize = (5, 5), dpi = 150)
sns.heatmap(cm, annot = True, ax = ax, cbar = False, fmt = 'd');
```



**Compute the accuracy, precision and recall.**

```
In [14]: from sklearn.metrics import recall_score

acc = accuracy_score(y_perov, y_predict)
prec = precision_score(y_perov, y_predict)
rec = recall_score(y_perov, y_predict)

print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(prec))
print('Recall: {}'.format(rec))
```

```
Accuracy: 0.8350694444444444
Precision: 0.8244047619047619
Recall: 0.8849840255591054
```


## 6745 Only: Customizing non-linear boundaries

In this problem, you will create a single custom feature that improves the separation performance as much as possible.

**Plot the  $y_{\text{perov}}$  as a function of  $r_A$  (Ang) and  $r_B$  (Ang) .**

```
In [15]: rA = X_perov[:, 3]
rB = X_perov[:, 4]

fig, ax = plt.subplots(figsize = (5, 5), dpi = 150)
ax.scatter(rA, rB, c = y_perov, alpha = .4)
ax.set_xlabel('rA')
ax.set_ylabel('rB');
```



**Build a baseline model based on logistic regression.**

Report the accuracy and precision of the baseline model.

```
In [16]: logreg = LogisticRegression(penalty = 'none')

logreg.fit(X_perov[:, [3, 4]], y_perov)
y_predict = logreg.predict(X_perov[:, [3, 4]])
```

```
acc = accuracy_score(y_perov, y_predict)
prec = precision_score(y_perov, y_predict)
```

```
print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(prec))
```

Accuracy: 0.6631944444444444


Precision: 0.6263269639065817

### Plot the prediction of the baseline model.

```
In [17]: fig, axes = plt.subplots(1, 2, figsize = (10, 4), dpi = 150)
```

```
axes[0].scatter(rA, rB, c = y_perov, alpha = .4)
axes[0].set_xlabel('rA')
axes[0].set_ylabel('rB')
```

```
axes[1].scatter(rA, rB, c = y_predict, alpha = .4)
axes[1].set_xlabel('rA')
axes[1].set_ylabel('rB');
```




**Create a new feature based on a non-linear combination of rA (Ang) and rB (Ang) .**

Plot the new feature as a function of rA (Ang) .



```
In [18]: X_new = np.exp((rA - 1.2)**2 + (rB - 0.6)**2)**(-6.9)
X_new = X_new.reshape(-1, 1)

fig, ax = plt.subplots(figsize = (5, 5), dpi = 150)
ax.scatter(rA, X_new, alpha = .4, c = y_perov)
ax.set_xlabel('rA')
ax.set_ylabel('new feature');
```



**Build a new model that includes rA (Ang) , rB (Ang) and your new feature.**

Report the accuracy and precision.

```
In [19]: X_new_matrix = np.append(X_perov[:, [3, 4]], X_new, 1)

logreg.fit(X_new_matrix, y_perov)
y_predict = logreg.predict(X_new_matrix)

acc = accuracy_score(y_perov, y_predict)
prec = precision_score(y_perov, y_predict)

print('Accuracy: {}'.format(acc))
print('Precision: {}'.format(prec))
```


Accuracy: 0.8038194444444444  
Precision: 0.8597122302158273

**Plot the result of your new model.**

```
In [20]: fig, axes = plt.subplots(1, 2, figsize = (10, 4), dpi = 150)

axes[0].scatter(rA, rB, c = y_perov, alpha = .4)
axes[0].set_xlabel('rA')
axes[0].set_ylabel('rB')

axes[1].scatter(rA, rB, c = y_predict, alpha = .4)
axes[1].set_xlabel('rA')
axes[1].set_ylabel('rB');
```



**Briefly explain how you decided on the feature.**

Type *Markdown* and LaTeX:  $\alpha^2$

### 3. Comparison of Classification Model

In this problem, you will compare the classification performance of three different models using the perovskite dataset.

**Choose three different classification models and import them.**

These could be models discussed in the lectures, or others that you have learned

about elsewhere.

```
In [21]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

rf = RandomForestClassifier()
logreg = LogisticRegression(solver = 'saga')
knn = KNeighborsClassifier()
```

### Make a hyperparameter grid for each model.

You should optimize at least one hyperparameter for each model.

```
In [22]: param_rf = {'n_estimators': [50, 100, 150], 'max_depth': [2, 3, 4]}
param_logreg = {'penalty': ['l1', 'l2', 'none']}
param_knn = {'n_neighbors': [2, 3, 4, 5]}
```

### Optimize hyperparameters.

First, you select a validation set using hold-out ( `train_test_split` ). Optimize hyperparameters using `GridSearchCV` on the training set.

```
In [23]: from sklearn.model_selection import GridSearchCV
import warnings

warnings.simplefilter('ignore')

X_train, X_test, y_train, y_test = train_test_split(X_perov, y_perov, test_size = .25)

rf_search = GridSearchCV(rf, param_rf, cv = 3)
logreg_search = GridSearchCV(logreg, param_logreg, cv = 3)
knn_search = GridSearchCV(knn, param_knn, cv = 3)

rf_search.fit(X_train, y_train)
logreg_search.fit(X_train, y_train)
knn_search.fit(X_train, y_train)

rf_best = rf_search.best_estimator_
logreg_best = logreg_search.best_estimator_
knn_best = knn_search.best_estimator_
```

### Compare the accuracy by predicting the results of the validation set.

```
In [24]: print('Accuracy of Random Forest: {}'.format(accuracy_score(y_test, rf_best.predict(X_test))))
print('Accuracy of Logistic Regression: {}'.format(accuracy_score(y_test, logreg_best.predict(X_test))))
print('Accuracy of k-Nearest Neighbors: {}'.format(accuracy_score(y_test, knn_best.predict(X_test))))

Accuracy of Random Forest: 0.9166666666666666
Accuracy of Logistic Regression: 0.6597222222222222
Accuracy of k-Nearest Neighbors: 0.9444444444444444
```

**Briefly describe your conclusions based on the results.**

Type *Markdown* and LaTeX:  $\alpha^2$